

Trevor Stahl

Project 6

CUDA Monte Carlo

CS 475e

Professor Bailey

I ran the code on my own personal machine. I have Windows 10 up to date as of 5/30/2019 with an AMD 2700x, 8 core hyper threaded CPU at 4.1 Ghz right now. The GPU is an Nvidia 1080ti with a mediocre stock cooler. 16.3 GBs of RAM. Background CPU process kept below 5% and background GPU kept below 10%.


I am using Visual studio 2017. I needed to retarget the solution and change the SDK version.

“GPU Device 0: "GeForce GTX 1080 Ti" with compute capability 6.1”


Table:

|    | 16000  | 32000   | 64000   | 128000  | 256000  | 512000  |
|----|--------|---------|---------|---------|---------|---------|
| 16 | 504.03 | 743.49  | 905.8   | 984.25  | 1036.81 | 1064.11 |
| 32 | 781.25 | 1041.67 | 1453.49 | 1785.71 | 1953.13 | 2066.12 |
| 64 | 822.37 | 1358.7  | 1459.85 | 1781.74 | 1968.5  | 2083.33 |

Graph of performance vs. NUMTRIALS with multiple curves of BLOCKSIZE:



Graph of performance vs. BLOCKSIZE with multiple curves of NUMTRIALS:



Patterns seen:

- A block size of 16 is consistently much worse than 32 or 64
- A block size of 64 performs significantly better than 32 only at NUMTRIALS = 32000
- A block size of 64 performs very similarly compared to a block size of 32 at all of the other points
- A NUMTRIALS of 32000 seems to produce a nearly straight line across 16, 32, 64 BLOCKSIZE values
- There is occasionally a very slight decrease in performance when holding NUMTRIALS constant but going from 32 to 64 blocksize. Most times 32 to 64 provides an increase

Why might these patterns be:

- I am not at all sure why 64 BLOCKSIZE only performs significantly better than 32 at a low NUMTRIALS value once
- Because a block size of 64 and 32 both provide warp sizes of 32 threads, a block size of 64 and 32 would perform similarly. (I think)
- A NUMTRIALS of 32000 probably has a straight line only by chance and the graph dimensions
- A block size of 64 might perform slightly worse because it might lead to having 2 thread blocks commute unnecessarily compared to 32 block size (Maybe)

Why is a BLOCKSIZE of 16 so much worse?

- A block size of 16 might be much worse because it does not provide for the ideal warp size and causes a doubling of the need for grids/blocks

What this means for proper use of GPU parallel computing:

- This all suggests to me that one needs to use a block size evenly divisible by 32, including 32. Again, I think this is because of the warps.

Probability:

The average probability for all of the different runs was 41.89%